

# Gandhar Kamat

(404) 729-0554 | gkamat1@student.gsu.edu | linkedin.com/in/gandhar-kamat | github.com/gandhar

## SUMMARY

Enterprise Architect with over 5 years of leadership experience in designing resilient mission-critical applications in AWS. Skilled in cloud computing, information management, application development, and virtualization. Experienced in communicating across organizational levels and with cross-functional teams to drive shared vision and foster culture of excellence.

## EXPERIENCE

**Software Developer SDE-II (Platform Engineering)**      **Dream11 Fantasy, Mumbai, India**      **Aug 2018 - Jul 2019**

- **Kubernetes Driven Development Solution** (EKS, EC2, RDS, ELB, Jenkins, Helm)
  - Designed and Implemented a cloud-agnostic solution for spawning complete or modular replicas of the production environment in Kubernetes.
  - Led development team of 7 engineers.
  - Designed the multi-stage docker images for use across environments.
  - Conducted hands-on training for 50+ developers on Docker and Kubernetes.
  - Explored various technologies for the containerization of applications and databases and other tools.
  - Developed scalable terraform modules for deploying Apache Spark, Kafka, RDS, etc. in AWS and Kubernetes.
- **Automated Apache Spark Scaling** ( Lambda, Spot Instances, Cloudwatch )
  - Designed a system to scale Apache Spark clusters on EC2 by predicting the expected traffic to the application.
  - Led to a 40% reduction in the cost of infrastructure of spark.
  - Used Cloudwatch to trigger lambda functions, which would call AWS APIs(python) to scale up and down according to projected capacity.
- **Anomaly Detection and Prevention** ( AWS WAF, CloudFront, Lambda@Edge)
  - Established Lambda@Edge processes to detect anomalous user behavior at the CDN level to proactively block potentially harmful requests into the infrastructure.
  - Automated CI-CD workflows for AWS Cloudfront and WAF deployments for streamlining development.
- **Real-Time Log Analysis** ( ElasticSearch, Lambda, SNS, Kinesis Firehose)
  - Created Elasticsearch-Fluentd-Kibana(EFK) for analyzing real-time application and AWS WAF logs.
  - Created actionable alerts for Site Reliability and On-Call teams( OpsGenie, NewRelic, ElastAlert )

**DevOps Engineer**      **LogiNext Solutions, Mumbai, India**      **Jul 2016 - Jan 2018**

- **Monolith to Microservices Migration** ( AutoScaling Groups, Netflix Eureka, Java )
  - Architected the migration of the production environment from a monolith to 40 highly available microservices using AutoScaling groups.
  - Explored frameworks for migration considering compatibility with the existing codebase.
  - Maintained cost after splitting into separate servers for high availability.
- **Tracking Service Development** ( Django, Elasticsearch, MongoDB, MySQL, Python )
  - Led the development of python microservices to generate alerts based using tracking information of delivery executives and long haul drivers.
  - Developed ETL pipelines for transforming time series location data from MySQL and MongoDB to Elasticsearch.
  - Introduced Database Indices and replication for distribution of load and performance improvement.
- **Infrastructure Automation** ( Ansible, Jenkins, Fabric )
  - Developed configuration management code to create production-like environments on multiple cloud providers( AWS, IBM Blumix, Alibaba Cloud ) as per business requirements.
  - Introduced automated Blue-Green deployments eliminating downtime during deployments and easy rollback.

**Software Engineer**      **Tesco Technology, Bangalore, India**      **Jun 2014 - Jun 2016**

- **PCI-DSS Compliant Deployments** ( Jenkins, Chef, CA Release Automation )
  - Designed a PCI-DSS compliant system for deploying the Payments infrastructure of Tesco.
  - Created Database deployment solution using featuring automated rollbacks on failures.

## CORE COMPETENCIES

---

- **Cloud:** AWS(EC2, S3, RDS, VPC, Lambda, CloudFront, SNS ), GCP, Alibaba Cloud, IBM Cloud
- **Programming:** Python (Numpy, Pandas, Scikit-learn), R, SQL, Java, Bash
- **Tools:** Tableau, Selenium, Git, Kubernetes, Docker, Ansible, Terraform, Jenkins, Fabric, NGINX, OpenVPN
- **Observability:** New Relic, Prometheus, Zabbix, Kibana, Fluentd, Logstash, Cloudwatch
- **Databases:** MySQL, MongoDB, Elasticsearch, Redis, Apache Spark, Apache Kafka

## EDUCATION

---

**J. Mack Robinson College of Business, Georgia State University** **Atlanta, GA, USA**  
MS in Data Science and Analytics December 2020

**Graduate Research Assistant** **Starr Insurance, Atlanta, GA, USA** **Jan 2020 - Now**

- Developing models to flag potentially high-value insurance claims in the first 6 months of the claim lifecycle.
- Created features from **unstructured data** by exploring accident descriptions and phone call transcripts.
- Working with time series transactional data, deriving insights and performing feature engineering. (PowerBI for viz)

**Graduate Research Assistant** **Andrew W. Mellon Foundation** **Aug 2019 - Dec 2019**

- Developed machine learning models to predict the likelihood of undergrad students to pursue higher degrees.
- Scraped relevant data from LinkedIn using **Selenium** to use as a supplemental data set.
- Achieved **97% precision and 83% sensitivity** using **Voting Classifier** with **XGBoost, Lasso, Random Forest**.

**National Institute of Technology** **Goa, India**  
B. Tech. in Electronics and Communication Engineering June 2014

## PROJECTS

---

- Performed **Sentiment Analysis** and **Topic Modelling** on movie review data, achieving **94% accuracy** in Binary Classification.
- Placed Top 10% in a Kaggle competition stacking **LightGBM, Support Vector Regression**, and **XGBoost** models.